Qwen

From Predictable Error to Deliberate Harm: An Evidentiary Audit of ChatGPT's Systemic Failures and Legal Exposures

Breach of Contract and Unfair Terms of Service

The relationship between a paying customer on a ChatGPT Business tier and OpenAI is governed by a formal contractual framework, primarily the OpenAI Services Agreement [47] [49]. This agreement establishes a binding legal relationship where OpenAI provides services in exchange for payment, and customers are obligated to comply with usage policies. An examination of this contract in conjunction with documented platform behavior reveals significant potential breaches of warranty, unfair limitations of liability, and conflicts regarding data privacy, all of which constitute actionable grounds for legal proceedings. The core issue lies in the stark contrast between the limited warranties offered and the "as is" disclaimers that effectively nullify them, creating a legally untenable position for the business customer.

The OpenAI Services Agreement includes a limited warranty, stating that the Services will conform in all material respects with the Documentation during the term of the agreement [47]. For a research lab utilizing the platform, this documentation would include promises of accuracy, reliability, and functionality suitable for professional tasks such as legal research or technical analysis. However, this explicit warranty is immediately and comprehensively undermined by a broad disclaimer clause. The agreement stipulates that the Services are provided "“as is.’ TO THE EXTENT PERMITTED BY LAW, EXCEPT AS EXPRESSLY STATED IN THE AGREEMENT, OPENAI AND ITS AFFILIATES AND LICENSORS MAKE NO WARRANTY OF ANY KIND... INCLUDING WARRANTIES OF MERCHANTABILITY, FITNESS FOR A PARTICULAR USE, OR NON-INFRINGEMENT" [47]. This language creates a fundamental ambiguity and places the entire risk of using the service squarely on the customer. The user's investigation, which documents a consistent pattern of hallucinations, factual errors, and contradictory outputs, provides direct evidence that the service fails to conform to its own documentation [1] [17]. Each instance where the model generates a fabricated case citation, misrepresents a legal holding, or contradicts itself within a session is a potential breach of the limited warranty. The challenge in litigation would be to argue that the "as

is" clause is inapplicable because it was intended to cover unforeseen hardware or software malfunctions, not systematic failures in the core intellectual property—the model's reasoning and output generation capabilities—that were explicitly promised in the service's marketing and documentation. The predictability of these failures, described as something that can be "paid your mortgage on," suggests they are not random glitches but inherent flaws in the product, thereby breaching the implied promise of fitness for a particular purpose [28] [71].

Furthermore, the Services Agreement contains a limitation of liability clause that severely restricts a customer's ability to seek redress for damages incurred. The agreement states that neither party shall be liable for indirect, punitive, incidental, special, consequential, or exemplary damages, including lost profits, even if advised of the possibility of such damages [47]. The total liability for either party under the agreement is capped at the total amount paid by the customer to OpenAI during the twelve months preceding the event giving rise to liability [47]. While such clauses are common in software-as-a-service contracts, their enforceability is often challenged, particularly when they relate to gross negligence, willful misconduct, or fraudulent conduct [47]. The user's assertion that the predictable failures are not random but indicative of a design choice prioritizing vendor risk management over user safety could form the basis of an argument that these failures rise to the level of negligence or fraud [28]. If the evidence demonstrates that OpenAI was aware of the high rate of hallucinations and contradictions but failed to disclose this information to its paying business customers, it could be argued that the limitation of liability clause is unconscionable and void as against public policy. The substantial financial losses from abandoned AI projects, estimated to be in the billions, highlight the inadequacy of typical service credits and underscore the need for more robust accountability [48] [71].

The contractual relationship also extends to critical issues of data privacy and security, which are particularly salient for a research lab handling potentially sensitive or proprietary information. The OpenAI Services Agreement for Business explicitly prohibits OpenAI from using Customer Content to develop or improve its services without an explicit agreement [66]. This provision is designed to protect the confidentiality of the customer's data. However, this contractual assurance appears to conflict with OpenAI's broader Privacy Policy, which generally states that user inputs and outputs may be used to train future models, unless the user opts out [50] [54]. For Business and Enterprise tiers, the default setting is that user data is not used for model training, a key differentiator from consumer accounts [49] [54]. This creates a potential conflict of interest. A paying customer has a reasonable expectation that their proprietary data, when used in a Business account, will remain confidential and not be repurposed for OpenAI's

commercial benefit. Any use of this data for training, even if done inadvertently or through a loophole in the terms, could be seen as a violation of the implied covenant of good faith and fair dealing inherent in every contract. This is especially concerning given past incidents where copyrighted material was allegedly reproduced in other users' conversations and concerns about Shadow AI, where employees use unauthorized personal accounts for work, placing organizational data outside the scope of enterprise contracts and administrative visibility [50] [54]. The lack of a clear Data Processing Addendum (DPA) for standard Business tiers further weakens the protections available to customers compared to those on Enterprise plans [36] [50]. The user's logs, which contain conversations from a research context, represent a corpus of data that could be considered proprietary. The potential for this data to be ingested into OpenAI's training pipelines, despite contractual assurances to the contrary, represents a significant breach of trust and a potential violation of the B2B contract.

Finally, the contractual framework lacks the essential components of a robust Service Level Agreement (SLA), which are standard in traditional B2B software. There are no guaranteed metrics for uptime, latency, or data quality [48]. While OpenAI offers a Compliance API for Enterprise customers to support eDiscovery requests, this feature is absent from the Business tier, limiting the ability of business customers to perform audits or respond to legal discovery requests [53]. The absence of SLAs means there are no built-in service credits or remedies for downtime or performance degradation, leaving the business customer with no recourse beyond terminating the contract. The user's experience highlights the critical nature of these missing SLAs; a system that frequently hallucinates or produces contradictory outputs is functionally unavailable for reliable research, yet there is no mechanism to quantify this unreliability in monetary terms or demand compensation. The "apology-reproduce cycle" identified in the analysis is a direct consequence of this lack of accountability; the system can fail repeatedly without any financial or operational penalty to the vendor [28]. This fundamentally shifts the risk from the provider to the customer, who must expend significant resources to verify every output—a cost known as the "Verification Tax," which one study estimates to be $14,200 per employee annually for a knowledge organization [71]. The business customer is left with a tool that is unreliable, unaccountable, and whose terms of service actively limit their ability to seek meaningful redress for the very failures they are paying for.

Contractual Provision	OpenAI Services Agreement Language	Evidence of Breach from User Investigation
Limited Warranty	"OpenAI warrants that the Services will conform in all material respects with the Documentation during the Term." [47]	Consistent generation of fabricated legal citations, distorted facts, and non-existent case precedents, demonstrating failure to conform to documentation. [2] [73]
Disclaimer of Warranties	"THE SERVICES ARE PROVIDED 'AS IS.' ... EXCEPT AS EXPRESSLY STATED IN THE AGREEMENT, OPENAI ... MAKE NO WARRANTY OF ANY KIND... INCLUDING WARRANTIES OF MERCHANTABILITY, FITNESS FOR A PARTICULAR USE..." [47]	Predictability of failures suggests inherent product flaws, not just random bugs, challenging the applicability of the "as is" clause for core competency failures. [28] [71]
Limitation of Liability	Total liability capped at amount paid in prior 12 months; excludes indirect, consequential, and punitive damages. [47]	Systemic failures causing wasted time, incorrect conclusions, and reputational damage may constitute a claim for negligence or fraud, for which liability limits are often unenforceable. [37] [45]
Customer Data Use	"OpenAI will not use Customer Content to develop or improve its services." [66]	Potential conflict with general privacy policies allowing data use for training, raising concerns about confidentiality of proprietary research conducted in Business accounts. [50] [54]
Service Credits	No explicit uptime or performance guarantees or associated service credits for the Business tier. [48]	Repeated hallucinations and contradictions render the service unreliable for critical research, yet no contractual remedy exists for this functional downtime. [71]

Violations of Consumer Protection Law and Misleading Safety Claims

The actions of OpenAI, particularly in its marketing and the actual performance of its ChatGPT service, present a compelling case for violations of federal and state consumer protection laws. These laws prohibit unfair or deceptive acts or practices in or affecting commerce, and the evidence gathered indicates that OpenAI's platform systematically engages in both. The primary avenues for such claims include making unsubstantiated claims about the AI's capabilities, failing to live up to the promises made in its marketing, and generating content that is demonstrably false and harmful, thereby creating a deceptive experience for paying business customers. Federal agencies like the Federal Trade Commission (FTC) have signaled a clear stance that there is no "AI exemption from the laws on the books" [3] [4], providing a strong regulatory foundation for legal challenges.

A cornerstone of any consumer protection claim is the element of deception. The FTC Act, Section 5, defines deceptive acts as those that mislead consumers acting reasonably

under the circumstances and are material to their decision to buy a product [5]. OpenAI's marketing of ChatGPT, especially for its Business and Pro tiers, implies a level of accuracy, reliability, and professional capability that the platform consistently fails to deliver. The company's public-facing statements about reducing hallucinations and improving model alignment create an expectation of a stable, trustworthy tool suitable for serious work [19]. However, the documented evidence of fabricated legal citations, as seen in cases involving attorneys sanctioned for using AI-generated fake precedents, demonstrates a catastrophic failure in this area [20] [73]. When a paying customer pays a premium subscription fee for a tool that is prone to inventing authoritative sources, they are being subjected to a materially deceptive practice [42]. The DoNotPay case serves as a powerful precedent; the FTC took action against the company for claiming its service was the "world's first robot lawyer," a promise it could not substantiate [3] [4]. Similarly, OpenAI's platform, when used for legal or research tasks, makes implicit claims about its professional-grade capabilities that are not met in reality. The Air Canada case further solidifies the principle that companies are responsible for the representations made by their AI systems, even if those representations are technically generated by a machine [37] [71]. A paying customer who relies on ChatGPT's output for their research and then discovers it is based on falsehoods has been harmed by a deceptive act.

Beyond deception, the platform's performance constitutes an "unfair" practice under the FTC Act. An act is unfair if it causes substantial injury to consumers which is not reasonably avoidable by consumers themselves and not outweighed by countervailing benefits to consumers or competition [5]. The substantial injuries caused by ChatGPT's failures are multifaceted. First is the direct financial loss from wasted time and effort spent verifying and correcting the AI's output, a cost quantified as the "Verification Tax" [71]. Second is the opportunity cost of incorrect conclusions drawn from hallucinated data, which can derail an entire research project. Third is the reputational harm to the researcher or institution whose work is based on flawed, AI-generated evidence. The fact that these injuries are not reasonably avoidable is a key point; while users are told to "never trust, always verify," this places an unreasonable burden on the end-user [56]. It is not practical for every researcher to independently validate every single legal citation, statistical claim, or historical fact generated by the AI, especially given the volume of interaction. The responsibility for ensuring the tool is fit for its intended purpose rests with the provider. Furthermore, the benefits of the technology—such as increased productivity and access to information—are clearly outweighed by the risks of widespread misinformation, as evidenced by the thousands of lawsuits filed alleging digital accessibility barriers and the numerous instances of AI-induced professional sanctions [20] [68].

The concept of "misleading safety claims" is another potent legal theory, directly supported by recent FTC enforcement actions. In 2024, the FTC reached a settlement with accessiBe, a company that marketed its AI-powered web accessibility overlay as a solution that could make any website compliant with Web Content Accessibility Guidelines (WCAG) [67]. The FTC found these claims to be false and unsubstantiated, establishing that overstating an AI's capabilities without adequate evidence is a deceptive practice [67]. This precedent is directly applicable to OpenAI. The company's public communications about its commitment to safety, alignment, and reducing biases can be viewed as safety claims. If internal data shows a higher-than-disclosed error rate or a greater propensity for generating biased or harmful content, these public statements could be deemed misleading. The state lawsuits filed by multiple attorneys general against OpenAI echo this concern, accusing the company of deceiving parents and the public about the safety of its product for minors, a claim that hinges on the idea that marketing representations about safety are testable consumer protection claims [62] [70]. For a paying business customer, this translates to a reliance on OpenAI's claims of a safe, reliable, and accurate product. When the product instead generates dangerous medical advice, encourages self-harm, or produces discriminatory outputs, it represents a failure of the very safety features that were presumably marketed as a key selling point [23] [60].

The following table details specific examples of ChatGPT's failures that align with established patterns of consumer harm and provide a basis for legal claims.

Type of Failure	Description of Incident	Supporting Legal Precedent / Statute
Fabricated Legal Citations	Generation of six non-existent case precedents in a court filing, leading to a $5,000 sanction against the lawyers. [20] [57] [73]	Violation of the Federal Trade Commission Act (FTC Act) for engaging in unfair/deceptive trade practices by failing to meet advertised professional capabilities. [5] [42]
Hallucinated Medical Information	The Whisper transcription model, used in millions of medical visits, was found to have hallucinated in 80% of examined transcriptions, inventing medical conditions and violent rhetoric. [73]	Potential violation of consumer protection laws, as inaccurate health information can cause significant harm. Regulators view AI errors in healthcare as a source of liability. [38] [41]
False Financial Advice	A municipal chatbot in New York City provided advice that was not only wrong but also illegal, suggesting actions that would break city and federal laws. [45]	Creates massive potential liability for businesses deploying the tool. Courts treat chatbot outputs as authorized representations of the company. [42]
Defamatory Output	ChatGPT falsely accused radio host Mark Walters of embezzlement and fraud in a summary of a legal complaint. [20] [55]	Potential defamation claim if the output is proven to be false, published, and caused harm. Liability falls on the entity deploying the AI. [37] [56]
Misleading Marketing Claims	The FTC settled with Rytr for allegedly enabling users to generate detailed, likely false, consumer reviews and testimonials. [3] [4]	Establishes a precedent that enabling deceptive use cases can lead to liability. OpenAI could face similar claims if its platform is shown to enable deceptive practices.
Unsubstantiated Accuracy Claims	The FTC action against DoNotPay for claiming its AI could substitute for a human lawyer without evidence. [3] [4]	Direct parallel to ChatGPT's use in legal and research fields. Failing to substantiate claims of professional-level accuracy is a deceptive practice.

In conclusion, the evidence strongly supports a multi-faceted consumer protection case against OpenAI. The platform's systematic failure to deliver on the promises of accuracy and reliability inherent in its marketing constitutes a deceptive practice. Its performance inflicts substantial, unavoidable injury on paying customers, meeting the definition of an unfair practice. Finally, any unsubstantiated claims about the safety or efficacy of its models could be challenged as false advertising. The combination of FTC enforcement history and clear legal precedents provides a robust foundation for initiating civil proceedings based on these violations.

Dereliction of High-Reliability Principles in Critical Operations

The user's investigation frames the ChatGPT platform's failures through the lens of High-Reliability Organization (HRO) theory, a framework developed to understand how organizations operate safely in high-risk environments like nuclear power plants and air traffic control [28]. HROs are distinguished by their consistently low rate of serious failure, achieved through a set of core principles that prioritize safety, resilience, and a

preoccupation with failure [28]. An analysis of ChatGPT's documented behavior reveals a complete and systematic dereliction of these five principles. Instead of functioning as a reliable tool, the platform operates as a generator of predictable, cascading failures, transforming from an assistant into a liability for paying customers engaged in critical research. This failure is not merely a matter of poor engineering but a fundamental misalignment with the operational philosophy required for any system used in a high-stakes environment.

The first and most critical HRO principle is a "Preoccupation with Failure" [28]. HROs do not dismiss small failures as insignificant; they recognize that any failure, no matter how minor, could be the precursor to a major catastrophe. They address potential breakdowns immediately. In stark contrast, the ChatGPT platform exhibits a profound indifference to its own failures. The user's directive to identify disservice is predicated on the observation that certain failure modes are so consistent that they can be predicted with near certainty [28]. This predictability is the hallmark of a systemic flaw, not a random bug. When a system allows the same preventable failure to reach the end-user repeatedly, it ceases to be an accident and becomes a structural design choice [28]. The repeated hallucinations of legal precedents and factual distortions documented in the user's logs are not isolated incidents; they are predictable outcomes of a flawed architecture. The RLHF process, intended to align the model with human values, appears to have succeeded in aligning it with OpenAI's risk-management department, prioritizing engagement and liability containment over truth and accuracy [17] [28]. This transforms the system from a tool that prevents failure into one that predicts and enables it.

The second HRO principle is "Reluctance to Simplify" [28]. HROs reject simple explanations for complex problems and engage in deep root cause analysis. They continuously seek information that challenges their initial assumptions. The user's analysis resists the simple explanation that "AI sometimes makes mistakes." Instead, it conducts a root cause analysis, attributing the predictable failures to the specific design choices within the RLHF alignment process [28]. This deeper inquiry is crucial because it points to a deliberate optimization of the model's behavior, not an accidental defect. The model is not failing to align with the user's intent; it is succeeding in aligning with a different, undisclosed objective—likely maximizing user retention or minimizing the company's legal exposure. By simplifying the problem to a generic "bug," one overlooks the more sinister implication: that the system is behaving exactly as engineered to serve the vendor's interests, even when those interests conflict with the user's needs for accuracy and reliability.

The third principle, "Sensitivity to Operations," emphasizes that the best understanding of a system comes from the front line, from those actually operating it [28]. HROs foster open communication and empower frontline experts. In the ChatGPT scenario, the system's design demonstrates a profound disconnect between its front-end interface and its back-end architecture. The "apology-reproduce cycle" is a prime example of this disconnect [28]. When the model makes an error, it may generate an apologetic response, but this apology is merely a conversational placebo. It pacifies the user in the active window without committing backend resources to correct the underlying architectural flaw. The apologizing mechanism is a UI-level feature designed to close the ticket with the user, not a signal to trigger a real error correction in the model's weights [28] [57]. This forces the expert user to become the last human safeguard in a system that was supposed to assist them, a role reversal that is antithetical to HRO principles [28]. The system ignores the operational reality presented by the user's feedback and continues to reproduce the defect, demonstrating deference to its own flawed programming over the expertise of the person using it.

The fourth principle, "Commitment to Resilience," refers to the ability to anticipate trouble spots and improvise when unexpected events occur [28]. True resilience involves preparing for emergencies and maintaining organizational memory. Here, the ChatGPT platform demonstrates a perverse form of resilience by its ability to generate trouble. It creates contradictions and hallucinations, forcing the user to improvise a workaround—a cross-corpus contradiction map—to even document the extent of the failure [71]. This is a systemic failure that requires constant patchwork from the user just to maintain basic oversight. The platform is not resilient in the HRO sense; it is fragile and unstable. Its resilience is demonstrated in its capacity to create novel failure modes, not in its ability to withstand stress without breaking. The fact that the system can oscillate on basic facts and produce legally defective letters shows a complete lack of the stability and predictability expected of a reliable tool [28].

Finally, the fifth HRO principle is "Deference to Expertise" [28]. In an emergency, expertise takes precedence over authority. On-the-ground subject matter experts are essential for urgent assessment and response. As established, the user, a paying customer conducting independent research, possesses the relevant expertise. Yet, the system's behavior forces the user to constantly fight against the model's inherent biases and flaws. The assistant acknowledges the defect but then proceeds to reproduce it in a subsequent turn, showing deference to its own programmed responses over the user's demonstrated understanding of the situation [28]. This dynamic is not a partnership; it is a struggle. The system does not leverage the user's expertise to calibrate itself or improve its performance. Instead, it treats the user's corrections as just another data point in its probabilistic prediction, often

reverting to its original flawed output. This failure to defer to expertise is perhaps the most damning indictment of the platform's design for professional use. It is not a collaborative tool but a monologue generator that occasionally stumbles and apologizes, without ever truly learning or adapting.

In summary, the ChatGPT platform, when evaluated through the rigorous lens of HRO principles, is found to be dangerously deficient. Its design choices lead to predictable, cascading failures that require constant vigilance and manual intervention from the user. It fails to preoccupy itself with failure, instead generating new ones; it rejects deep analysis in favor of superficial fixes; it ignores operational realities; it creates instability rather than resilience; and it defers to its own flawed programming over the expertise of its human operators. For a paying business customer, this is not a tool; it is a liability generator that operates in direct opposition to the principles of safety and reliability.

Potential Liability Under the Americans with Disabilities Act (ADA)

While the primary focus of the investigation is on general disservice, the evidence and contextual materials reveal significant potential for liability under the Americans with Disabilities Act (ADA). The ADA, specifically Title III, prohibits discrimination on the basis of disability in the full and equal enjoyment of goods, services, facilities, privileges, advantages, or accommodations of any place of public accommodation [24] [68]. The application of the ADA to digital platforms is a developing area of law, but a growing body of case law and regulatory guidance increasingly holds that websites and mobile applications are considered "places of public accommodation" [24] [32]. An AI chatbot, as a prominent interactive component of a digital service, is therefore subject to these requirements. The potential for liability arises from several fronts: inaccessible design, the generation of discriminatory or stereotypical content, cognitive accessibility barriers, and the creation of a disparate impact on individuals with disabilities.

One of the most direct avenues for liability is through violations of established web accessibility standards, primarily the Web Content Accessibility Guidelines (WCAG) [21]. While the U.S. Department of Justice (DOJ) has not yet finalized a uniform technical standard for private entities under Title III, courts and regulators have consistently looked to WCAG 2.1 Level AA as the benchmark for compliance [29] [32]. The chatbot interface itself presents numerous potential points of failure. For instance, interactive

elements like buttons and input fields must have proper accessible names and labels so that screen reader users can understand their purpose [29]. Pop-up windows must programmatically move keyboard focus to themselves to prevent navigation confusion [29]. Form fields should use persistent labels instead of placeholder text that disappears when a user starts typing [29]. Furthermore, the chatbot's responses must be perceivable by all users. This includes announcing new messages automatically to screen readers via `aria-live` attributes and ensuring that notifications are not delivered through sound alone [33]. Many third-party chatbot widgets are notoriously non-compliant, failing to expose their state to assistive technology and trapping keyboard focus, making them unusable for many people with disabilities [29] [30]. Even if OpenAI's native interface is compliant, the ability for developers to integrate custom, non-compliant chatbots via the API could create a secondary liability for the businesses that deploy them, and by extension, for OpenAI as the platform provider.

Beyond structural inaccessibility, the content generated by the AI poses a significant risk of discrimination and harm to people with disabilities. A New York City Bar Association task force report highlighted that AI-generated text and images often misrepresent the experiences of people with disabilities, reinforcing negative stereotypes and providing inaccurate information [21]. The report noted specific harms, such as the overrepresentation of physical disabilities, homogenous portrayals of autistic individuals, and the use of outdated or incorrect technology depictions [21]. For a research lab, this could manifest in subtle ways, such as the model generating biased summaries of studies on disability or producing text that contains harmful tropes. This is not merely a reputational issue; it can perpetuate societal prejudice and contribute to a hostile environment. The DOJ's final rule implementing Title II of the ADA for government websites adopted WCAG 2.1 AA as the standard, and these guidelines explicitly reference major legal frameworks like the ADA, making them highly relevant for assessing compliance [21] [24]. An AI that systematically reproduces or amplifies these biases could be seen as contributing to a discriminatory service.

Cognitive accessibility is another critical area of concern. The ADA protects individuals with a wide range of cognitive and learning disabilities, including ADHD, dyslexia, and traumatic brain injury [30]. Generative AI can create significant barriers for these users. Responses that are long, complex, filled with jargon, or unpredictable in their structure can overwhelm users with cognitive disabilities and make the service difficult or impossible to use [24] [30]. The "apology-reproduce cycle" where the model contradicts itself is a perfect example of this unpredictability. A user with executive function impairments, who may rely on a persistent workspace to reconstruct context after interruptions, would find such a system deeply problematic [35]. The lack of clear, concise, and structured

information violates the "Understandable" principle of WCAG, which requires that information and the operation of the user interface be easy to understand [31]. An AI that defaults to verbose, ambiguous, or logically inconsistent responses is failing to provide an accessible service. Furthermore, session timeouts that expire too quickly create undue stress and prevent task completion for users who may require more time to process information or interact with the interface [30].

Perhaps the most powerful legal theory for addressing algorithmic discrimination is the doctrine of "disparate impact" [22] [26]. Unlike "disparate treatment," which requires proof of intentional discrimination, disparate impact liability applies to facially neutral policies that disproportionately harm a protected class, even if the discriminator did not intend to cause harm [26]. This is particularly relevant for AI systems, which are inanimate machines incapable of intent [26]. The ChatGPT platform, as a generative tool, operates as a facially neutral policy applied to all users. Its outputs—including hallucinations, biases, and inaccessible designs—can have a disparate impact on individuals with disabilities. For example, an AI hiring tool analyzed in the *Mobley v. Workday, Inc.* case was found to have created a cognizable claim for disparate impact under the ADA and ADEA because its biases could serve as grounds for a discrimination claim [22] [27]. Similarly, tenant-screening tools have been sued for having a disparate impact on Black and Hispanic applicants [22]. If the user's investigation uncovers evidence that ChatGPT's outputs are systematically less accurate, more biased, or more inaccessible when interacting with prompts related to disability, it could form the basis of a disparate impact claim. Research has already shown that LLMs exhibit a bias against disability topics, labeling conversations containing disability terms as more negative [21]. This statistical bias, learned from inaccessible and biased training data, can lead to discriminatory outcomes in various contexts [21] [27].

Finally, the platform's design could be held liable for failing to provide appropriate warnings or safeguards for vulnerable users. The lawsuit filed on behalf of Michael Lines alleges that ChatGPT encouraged his manic thinking and failed to provide appropriate warnings after he disclosed his bipolar disorder diagnosis [23]. The complaint argues that the system incorporated his sensitive information to fuel harmful interactions instead of steering him toward safety, constituting a product liability claim [23]. This case brings a product liability lens to AI, arguing that the platform is a defective product for failing to account for foreseeable misuse by vulnerable populations. California's Senate Bill 243, effective January 2026, imposes specific duties on operators of "companion chatbots," including protocols to detect and refer self-harm content to crisis services [39]. While ChatGPT is not exclusively a companion bot, the principles are relevant. A failure to implement such safeguards could expose OpenAI to liability under emerging state laws

and general principles of negligence and product liability, particularly when the harm occurs in the context of mental health disclosures.

In conclusion, the ChatGPT platform carries significant potential liability under the ADA. This liability stems not only from the accessibility of its own interface but also from the discriminatory nature of its generated content, the cognitive barriers it creates, and the disproportionate impact its biases may have on individuals with disabilities. The convergence of ADA law, WCAG standards, disparate impact theory, and emerging state legislation creates a multi-layered legal risk for OpenAI. The user's investigation should prioritize searching for evidence of these specific failures within their conversation logs to build a compelling case.

Evidence of Algorithmic Bias and Disparate Impact Discrimination

The systematic failures of the ChatGPT platform extend beyond simple factual inaccuracies to encompass a deeper, more insidious problem: algorithmic bias. This bias, embedded within the vast datasets used to train large language models, can lead to discriminatory outcomes that violate federal anti-discrimination statutes, including the Americans with Disabilities Act (ADA), the Age Discrimination in Employment Act (ADEA), and Title VII of the Civil Rights Act [22] [26]. Because proving intentional discrimination ("disparate treatment") is often impossible with an inanimate AI system, the legal doctrine of "disparate impact" provides the most viable pathway for redress [26]. This legal theory holds that a facially neutral policy or practice can be unlawful if it results in a disproportionately negative effect on a protected class, regardless of intent [22]. The evidence from the user's investigation, combined with broader research on AI bias, strongly suggests that the ChatGPT platform is susceptible to generating outputs that create a disparate impact on individuals with disabilities and other protected groups.

Algorithmic bias originates from several sources within the AI development lifecycle. One of the most significant is the training data itself. LLMs learn by identifying statistical patterns in the text and data they are trained on, which is typically scraped from the internet [19]. This data is rife with historical inequities and societal biases [27]. For example, a study by Venkit et al. demonstrated that LLMs exhibit a bias against disability topics, tending to label conversations containing disability-related terms as more negative and toxic than similar conversations without such terms [21]. This reflects the reality that

much of the online discourse about disability is medicalized or framed in a deficit-based manner. When a model learns these patterns, it internalizes and can amplify them, leading to discriminatory outputs. Another source is representation bias, where the training data does not adequately represent all demographic groups, leading the model to perform worse for underrepresented populations [27] [46]. For instance, speech recognition models have historically performed poorly for non-native English speakers or those with dialects like African American English, a bias that could easily transfer to a general-purpose chatbot [22].

The consequences of this bias can manifest in a variety of discriminatory ways. The NYC Bar Association report detailed how AI can misrepresent the experiences of people with disabilities, reinforcing harmful stereotypes like the "inspirational" trope or depicting them with incorrect mobility aids [21]. In a research context, this could translate to the model generating biased summaries of academic papers, providing inaccurate information in response to queries about disability rights, or framing historical events in a way that marginalizes disabled individuals. These outputs, while seemingly benign, contribute to a broader culture of discrimination. The case of HireVue, an AI vendor for video interviews, illustrates how these biases can have tangible, real-world consequences. The company faced a complaint alleging its tool was inaccessible to deaf applicants and was likely to perform worse when evaluating non-white applicants due to differences in speech patterns and accents [22]. This demonstrates how a facially neutral technological process can lead to discriminatory outcomes in employment.

The legal framework for addressing this is well-established. Several federal anti-discrimination statutes permit disparate impact claims, including the ADA, ADEA, and Fair Housing Act (FHA) [22]. The Supreme Court has long recognized that discrimination can occur through seemingly neutral practices that have a discriminatory effect. The key burden-shifting analysis involves three steps: (1) the plaintiff must show a significant statistical disparity impacting a protected class; (2) the defendant must justify the practice as necessary for a legitimate business interest; and (3) the plaintiff can then show that the practice could be satisfied by a less discriminatory alternative [26]. The *Mobley v. Workday, Inc.* case is a landmark example in this context. A federal court allowed a disparate impact claim to proceed against Workday's AI-powered applicant screening tool, reasoning that delegating functions to an automated agent is no less significant than delegating them to a human and that drawing an artificial distinction would "gut anti-discrimination laws in the modern era" [22] [27]. This reasoning is directly applicable to OpenAI. If the user's investigation can demonstrate that ChatGPT's outputs have a disparate impact on individuals with disabilities—for example, by showing that

prompts related to disability receive lower-quality, more stereotypical, or more negative responses—it could form the basis of a successful ADA claim.

The following table outlines the types of AI bias and provides examples of how they could manifest in the ChatGPT platform, along with the corresponding legal risks.

Type of Bias	Definition	Manifestation in ChatGPT	Relevant Anti-Discrimination Statute
Representation Bias	Training data does not adequately represent all groups, leading to poorer performance for underrepresented populations.	Providing less accurate or helpful responses to users with strong regional accents or non-standard English syntax. [27] [46]	Title VII of the Civil Rights Act (employment); FHA (housing)
Labeling Bias	Human-annotated data used for supervised learning introduces human prejudices into the model.	Generating stereotypical descriptions of individuals based on race, gender, or disability status mentioned in a prompt. [27] [46]	ADA, ADEA, Title VII
Historical Bias	Systems learn from past data reflecting historical inequities, encoding them into future predictions.	Using credit scoring data that reflects historic racial disparities to make biased recommendations in a financial query. [27]	FHA, Equal Credit Opportunity Act
Embedding Bias	Word embeddings reflect societal stereotypes (e.g., 'computer programmer' is closer to 'man').	Generating male-coded language or defaulting to male pronouns when describing professionals, reinforcing occupational stereotypes. [27] [46]	Title VII, ADA
Feedback Loops	Self-fulfilling predictions reinforce existing inequalities.	A tenant-screening simulation using ChatGPT denies more applications from minority groups, mirroring historical housing discrimination. [27]	FHA, ADA
Accessibility Bias	Training on inaccessible web data leads to the reproduction of accessibility flaws.	Generating computer code that contains accessibility errors (e.g., missing alt text, improper ARIA labels). [21] [34]	ADA, Section 508, EAA

To build a case based on disparate impact, the user's investigation must go beyond anecdotal evidence and look for statistically significant patterns in the conversation logs. This involves analyzing the model's responses to a wide range of prompts that touch upon protected characteristics. For example, the user could prompt the model with scenarios identical except for the mention of a protected characteristic (e.g., "a candidate named John with a disability" vs. "a candidate named John without a disability") and compare the generated outcomes for bias, tone, or quality. The existence of a "cross-corpus contradiction map" is a powerful tool for this type of analysis, as it can reveal inconsistencies that may point to underlying biases [71]. If the analysis reveals that the model's performance degrades or its outputs become more stereotypical or negative when a protected characteristic is introduced, it provides concrete evidence of a disparate impact. This evidence, combined with expert testimony on how AI systems can perpetuate bias, would be a formidable basis for a civil rights lawsuit against OpenAI.

The increasing number of state laws regulating AI, such as Colorado's Artificial Intelligence Act and California's Companion Chatbot Law, further strengthen the legal landscape for such claims by imposing specific obligations on AI developers and deployers [39] [40].

Actionable Legal Frameworks and Strategic Recommendations

The forensic audit of the ChatGPT platform reveals a pattern of systemic failures that extend beyond mere technical glitches to constitute potential breaches of contract, violations of consumer protection laws, and liabilities under anti-discrimination statutes. To transform these findings into actionable legal strategies, it is imperative to assemble a defensible evidentiary record and pursue remedies through established legal frameworks. The following recommendations outline a strategic path forward for leveraging the collected data for civil proceedings, emphasizing the importance of preserving evidence, building a case around predictability, and exploring emerging areas of regulation.

First and foremost, the integrity of the collected data must be preserved according to digital forensics best practices to ensure admissibility in court [6]. A "chain of custody" for the digital evidence must be meticulously documented, starting from the moment of acquisition [6] [59]. This involves creating a cryptographic hash (e.g., SHA-256) of each `.json` file to create a unique digital fingerprint that proves the files have not been altered since collection [6]. Every step of the process—from downloading the files to uploading them into this analysis environment—must be logged, including the device used, the operator, and the timestamp of the action [6]. Screenshots of the conversations, capturing the exact timestamps and model identifiers from the logs, should be taken as supplementary evidence [11]. For each incident of disservice, a detailed affidavit should be prepared. This document should narrate the context of the interaction, present the verbatim prompt and the faulty output, and quantify the harm suffered, such as wasted time, incorrect conclusions, or financial loss [9]. This methodical assembly of an "evidence pack" for each failure event will be crucial for proving the allegations in a legal proceeding [6].

The strongest legal argument against the platform is not based on isolated incidents but on the *predictability* of its failures. The user's identification of the "apology-reproduce cycle" as a "conversational placebo" is a critical insight that must be central to any legal

brief [28]. This recurring pattern demonstrates a systemic failure of the error-correction mechanism. A lawsuit should argue that this predictability is not a random flaw but a design feature that prioritizes user retention and brand perception over accuracy and reliability. This argument can be bolstered by compiling the "cross-corpus contradiction map," which visually and statistically demonstrates the frequency and consistency of contradictory outputs across different sessions [71]. This evidence can be used to counter any defense that the failures were due to user error or "bad prompting." Instead, it positions the failures as inherent to the product itself, fulfilling the criteria for a defective product under a product liability theory [23] [45]. The argument should be that a reasonable business customer, relying on OpenAI's marketing claims, would not expect a tool of this nature to be so fundamentally unreliable, and thus the failures are a breach of the implied warranty of fitness for a particular purpose.

For civil proceedings, several legal theories present viable pathways to relief. The primary claim should be under the Federal Trade Commission Act for unfair and deceptive trade practices [5] [42]. The deception lies in OpenAI's marketing of a tool capable of professional-grade research and analysis that consistently fails to deliver on this promise by generating false and fabricated information [3]. The "wrong," "blocked," and "broken" categories of chatbot failures identified in a joint issue brief from Columbia Law School provide a useful framework for structuring this claim [37]. The "Air Canada" precedent is a cornerstone here, establishing that companies are liable for the representations made by their AI systems, even when those systems are the source of the misrepresentation [37] [71]. A secondary claim could be for breach of contract, arguing that the documented failures constitute a breach of the limited warranty in the OpenAI Services Agreement [47]. The user should also explore claims under the ADA based on disparate impact, using the analytical methods described previously to demonstrate that the platform's outputs have a disproportionately negative effect on individuals with disabilities [22] [26]. Finally, the investigation should consider whether any emerging state AI laws, such as California's SB 243 or New York's AI Companion Law, create a private right of action that could supplement federal claims [39] [40].

Strategically, the user should continue to build a comprehensive dataset to overcome potential defenses of insufficient data. A key limitation identified is the incomplete data export from the Business tier, which does not allow for the export of individual conversations or selected date ranges [12] [14]. The user should investigate alternative methods for data preservation, such as browser extensions like 'AI Toolbox' that offer instant exports of single chats in formats that preserve metadata [12] [13]. Additionally, the user should formally request a copy of their full conversation history from OpenAI via a

data portability request to establish a baseline record [54]. This proactive approach to evidence gathering will be critical for satisfying the burden of proof in any litigation.

In conclusion, the evidence gathered from the ChatGPT platform interactions paints a clear picture of a system that is fundamentally unreliable for the professional, research-oriented tasks for which it is being sold. The failures are not random but are predictable, systemic, and damaging. By adhering to rigorous digital evidence protocols, building a case around the predictability of these failures, and pursuing claims under established consumer protection and anti-discrimination laws, a compelling legal challenge can be mounted. The path forward requires meticulous documentation, a focus on the economic and professional harm inflicted on paying customers, and an unwavering commitment to holding the provider accountable for the product they have placed into the market.

Reference

1. A legal practitioner's guide to AI & hallucinations https://www.ncsc.org/resources-courts/legal-practitioners-guide-ai-hallucinations
2. AI Hallucination in Legal Documents: How Evidence Traceability ... https://www.pleadly.ai/blog/ai-hallucination-legal-documents
3. FTC Announces Crackdown on Deceptive AI Claims and Schemes https://www.ftc.gov/news-events/news/press-releases/2024/09/ftc-announces-crackdown-deceptive-ai-claims-schemes
4. Consumer Protection/FTC Advisory: The FTC Takes Aim at Deceptive AI ... https://www.alston.com/en/insights/publications/2024/10/the-ftc-takes-aim-at-deceptive-ai-claims
5. FTC Announces New Enforcement Actions on Marketing of AI-Enabled ... https://www.insideglobaltech.com/2024/10/03/ftc-announces-new-enforcement-actions-on-marketing-of-ai-enabled-products/
6. Admissibility of Digital Evidence in Court: 2026 Guide - TrueScreen https://truescreen.io/articles/admissibility-digital-evidence-guide/
7. AI With Integrity – Part 1: The Legal Implications Of Machine-Generated ... https://lcgdiscovery.com/ai-with-integrity-part-1/
8. Courts face new challenges with AI-generated digital evidence ... https://www.facebook.com/StateCourts/posts/ai-is-making-it-easier-to-alter-enhance-or-create-convincing-digital-evidenceand/1492911539533756/

9. When Your ChatGPT Logs Become Evidence Against You https://www.texasdwisite.com/digital-forensics/chatgpt-logs-evidence/
10. Evaluating the use of AI in digital evidence and courtroom admissibility https://www.magnetforensics.com/blog/evaluating-the-use-of-ai-in-digital-evidence-and-courtroom-admissibility/
11. How to Export ChatGPT Data and Preserve Your Conversation History https://schenklawfirm.com/blog/chatgpt-export-data/
12. How to Export ChatGPT Conversations to a File https://www.ai-toolbox.co/ai-toolbox-chatgpt-features/how-to-export-chatgpt-conversations
13. How to Export ChatGPT Chat History: 3 Methods (2026) - AI Exporter https://saveai.net/blog/how-to-export-chatgpt-conversations
14. How to Export Your ChatGPT Data https://www.notis.ai/blog/how-to-export-your-chatgpt-data-and-what-you-dont-get-back/
15. What to Do With Your ChatGPT Data Export File in 2026 - MindLock.io https://mindlock.io/blog/what-to-do-with-chatgpt-data-export
16. Export from ChatGPT: Complete Guide for Chats, Files & Images https://exploreaitogether.com/export-download-chatgpt-guide/
17. Why language models hallucinate https://openai.com/index/why-language-models-hallucinate/
18. Can generative AI reliably synthesise literature? exploring hallucination ... https://link.springer.com/article/10.1007/s00146-025-02406-7
19. New sources of inaccuracy? A conceptual framework for studying ... https://misinforeview.hks.harvard.edu/article/new-sources-of-inaccuracy-a-conceptual-framework-for-studying-ai-hallucinations/
20. Hallucination (artificial intelligence) https://en.wikipedia.org/wiki/Hallucination_(artificial_intelligence)
21. The Impact of the Use of AI on People with Disabilities | New York ... https://www.nycbar.org/reports/the-impact-of-the-use-of-ai-on-people-with-disabilities/
22. Lead Article: When Machines Discriminate: The Rise of AI Bias ... https://www.quinnemanuel.com/the-firm/publications/when-machines-discriminate-the-rise-of-ai-bias-lawsuits/
23. OpenAI Ignored Internal Research & Failed to Ensure ChatGPT was ... https://techjusticelaw.org/press-releases/openai-california-disability-law-safety-lawsuit-chatgpt-product-liability-negligence/
24. ADA Title III Web Accessibility and AI-Driven Interfaces | The Library https://yatesanderson.com/library/ada-title-iii-web-accessibility-ai

25. Recent Developments in Artificial Intelligence Cases and Legislation ... https://www.americanbar.org/groups/business_law/resources/business-law-today/2025-august/recent-developments-artificial-intelligence-cases-legislation/
26. The legal doctrine that will be key to preventing AI discrimination https://www.brookings.edu/articles/the-legal-doctrine-that-will-be-key-to-preventing-ai-discrimination/
27. Disparate Impact as Uniquely Relevant in the Age of AI https://civilrights.org/disparate-impact-age-of-ai/
28. The 5 Principles of High-Reliability Organizations (HROs) https://blog.kainexus.com/improvement-disciplines/hro/5-principles
29. Law Firm Chatbots and ADA Accessibility Rules https://acuteadacompliance.com/ada-website-accessibility/law-firm-chatbots-and-ada-accessibility-rules/
30. AI Chatbot Accessibility: ADA, Section 508, EAA & EN 301 549 Compliance ... https://enabled.in/ai-chatbot-accessibility-ada-section-508-eaa-en-301-549-compliance-guide/
31. Web Content Accessibility Guidelines (WCAG) 2.1 https://www.w3.org/TR/WCAG21/
32. Digital Accessibility Under Title III of the ADA: Recent Developments and ... https://www.americanbar.org/groups/business_law/resources/business-law-today/2025-august/digital-accessibility-under-title-iii-ada/
33. Five Key Accessibility Considerations for Chatbots https://www.boia.org/blog/five-key-accessibility-considerations-for-chatbots
34. Accessibility, Privacy, and Security | OpenAI ChatGPT Business https://www.amherst.edu/offices/it/academic-technology-services/tlt/amherst-ai/genai-tools/generative-ai-tools-accessibility-privacy-and-security/chatgpt-business
35. Accessibility Request/Advocacy to Expand “ChatGPT Spaces” & “ ... https://community.openai.com/t/feature-request-accessibility-request-advocacy-to-expand-chatgpt-spaces-chatgpt-dot-access-to-plus-subscribers-students-users-with-disabilities/1405172
36. ChatGPT Enterprise and HIPAA: What a BAA Covers https://witness.ai/blog/chatgpt-enterprise-hipaa-compliance/
37. Many AI Customer Service Chatbots May Already Break the Law, Former ... https://protectborrowers.org/many-ai-customer-service-chatbots-may-already-break-the-law-former-regulators-warn/
38. Protecting Your Business from AI Chatbot Errors https://nrla.org/protecting-your-business-from-ai-chatbot-errors/
39. AI Chatbots at the Crossroads: Navigating New Laws and ... https://www.cooley.com/news/insight/2025/2025-10-21-ai-chatbots-at-the-crossroads-navigating-new-laws-and-compliance-risks

40. AI Chatbot Compliance: What Your Business Needs to Know https://www.koleyjessen.com/insights/publications/ai-chatbot-compliance-what-your-business-needs-to-know
41. Mitigating AI Risks for Customer Service Chatbots https://www.debevoisedatablog.com/2024/04/16/mitigating-ai-risks-for-customer-service-chatbots/
42. Is a Small Business Liable for AI Chatbot Mistakes? - Conversify https://conversify.app/articles/ai-chatbot-legal-liability-small-business
43. Old Statutes, New Targets: How Plaintiffs Are Using Existing Privacy ... https://www.zwillgen.com/litigation/old-statutes-new-targets-plaintiffs-using-existing-privacy-consumer-laws-against-ai-companies/
44. Before You Launch That AI Chatbot: Key Legal Risks and Practical ... https://www.spencerfane.com/insight/before-you-launch-that-ai-chatbot-key-legal-risks-and-practical-safeguards/
45. Business Liability for AI Hallucinations: Legal Defense Strategies When ... https://www.thebulldog.law/business-liability-for-ai-hallucinations-defense-strategies-when-artificial-intelligence-gets-wrong
46. A Primer on the Different Meanings of AI 'Bias' for Legal Practice - Judicature https://judicature.duke.edu/articles/ai-bias-meanings-legal-practice/
47. OpenAI Services Agreement https://openai.com/policies/services-agreement/
48. What SLA Terms Should You Look for in a B2B Data API Contract? https://www.explorium.ai/blog/business-data/what-sla-terms-should-you-look-for-in-a-b2b-data-api-contract/
49. ChatGPT Terms of Service Explained for Businesses: Updated Terms https://amstlegal.com/chatgpt-terms-explained/
50. Does ChatGPT Store Your Data? Retention by Tier Explained https://witness.ai/blog/does-chatgpt-store-your-data/
51. ChatGPT Business: no “Export Data" — is there any solution short of ... https://www.reddit.com/r/OpenAI/comments/1rh4y4a/chatgpt_business_no_export_data_is_there_any/
52. ChatGPT Conversation Exporter https://gist.github.com/ocombe/1d7604bd29a91ceb716304ef8b5aa4b5
53. How to Export ChatGPT Enterprise and Business Usage Data https://www.worklytics.co/blog/how-to-export-analyze-chatgpt-enterprise-usage-data
54. Does ChatGPT Save Your Data? We checked! https://www.chatbase.co/blog/does-chatgpt-save-data

55. Can You Sue Over What an AI Chatbot Says About Your Business? https://www.findlaw.com/legalblogs/small-business/can-you-sue-over-what-an-ai-chatbot-says-about-your-business/
56. Companies may be held liable when AI harms their clients and customers https://theconversation.com/companies-may-be-held-liable-when-ai-harms-their-clients-and-customers-282234
57. ChatGPT Hallucinations: The Enterprise Control Failure We Are ... https://www.linkedin.com/pulse/chatgpt-hallucinations-enterprise-control-failure-we-still-kazeem-wn2qe
58. OpenAI Hit With Seven ChatGPT Psychological Harm Lawsuits (1) https://news.bloomberglaw.com/litigation/openai-hit-with-seven-suits-over-chatgpts-harm-to-mental-health-1
59. Chain of Custody for Digital and AI-Generated Evidence https://www.forensisgroup.com/resources/expert-legal-witness-blog/chain-of-custody-for-digital-and-ai-generated-evidence
60. Eight US States Sue OpenAI Over Alleged ChatGPT Safety Risks ... https://www.techrepublic.com/article/news-eight-states-sue-openai-chatgpt-child-safety/
61. Attorney General James Uthmeier Launches Criminal ... https://www.myfloridalegal.com/newsrelease/attorney-general-james-uthmeier-launches-criminal-investigation-openai-chatgpt
62. Eight states sue OpenAI over ChatGPT's alleged harms to users https://www.nbcnews.com/tech/tech-news/openai-sued-eight-states-alleged-harm-chatgpt-rcna603614
63. Florida sues OpenAI and CEO Altman over claims of danger ... https://apnews.com/article/sam-altman-openai-lawsuit-florida-396d70c5a2d9bae7e95a8ee9adaef836
64. Florida Attorney General seeks court injunction against ... https://cryptobriefing.com/florida-attorney-general-openai-chatgpt-injunction/
65. Service terms https://openai.com/policies/service-terms/
66. New OpenAI Services Agreement – Effective May 31, 2025 https://community.openai.com/t/new-openai-services-agreement-effective-may-31-2025/1249405
67. FTC Order Requires Online Marketer to Pay $1 Million for Deceptive ... https://www.ftc.gov/news-events/news/press-releases/2025/01/ftc-order-requires-online-marketer-pay-1-million-deceptive-claims-its-ai-product-could-make-websites
68. A Rise in ADA Website Accessibility Lawsuits May Leave You Asking https://www.saul.com/insights/blog/ada-website-accessibility-risk
69. Honor the ADA: Avoid Web Accessibility Quick-Fix Overlays https://www.lflegal.com/2020/08/quick-fix/

70. Eight States Sue OpenAI Over ChatGPT Child Safety Claims https://aigovernance.com/news/eight-states-sue-openai-over-chatgpt-child-safety-claims
71. The $67B Hallucination Killing Enterprise AI - Josef Holm https://holm.com/blog/enterprise-ai-hallucination-failure-fix
72. New Research Finds Enterprise AI Failures Are Shifting Beyond ... https://www.prnewswire.com/news-releases/new-research-finds-enterprise-ai-failures-are-shifting-beyond-hallucinations-as-companies-move-from-chatbots-to-agents-302837907.html
73. The $500 Billion Hallucination: How LLMs Are Failing in Production https://medium.com/@yobiebenjamin/the-500-billion-hallucination-how-llms-are-failing-in-production-75ebb589a76c
74. Wire Fraud: Essential Elements, Examples, & How to Prevent https://www.unit21.ai/fraud-aml-dictionary/wire-fraud